

Scala Performance Considerations



Nermin Serifovic

ferrari48.com

About me

- ~ 10 years of experience building Java enterprise apps
- currently a Principal Software Engineer at Constant Contact
- Scala enthusiast since 2009
- co-founded Boston Area Scala Enthusiasts

@higherkinded

Methodology

Stopwatch benchmarking

Followed usual micro-benchmark writing rules:

- Warm-up phase
- Test subroutine outside of `main` method
- Series of measurements

HotSpot VM in server mode

Scala optimise flag made no difference

Scala performance

“Run-time performance is *usually* on par with Java programs”



Tight loops - Java

```
int s = 0

for (int i = 1; i <= 2000; i++)
    for (int j = 1; j <= 2000; j++)
        for (int k = 1; k <= 2000; k++)
            s += 1
```

90 μ s

Tight loops - Scala


```
var s = 0

for (i <- 1 to 2000;
     j <- 1 to 2000;
     k <- 1 to 2000)
  s += 1
```

1.3 s


```
(1 to 2000).foreach(  
  i => (1 to 2000).foreach(  
    j => (1 to 2000).foreach(k => s += 1)
```

```
k => s += 1
```



```
new Function1[Int, Unit]  
{  
  def apply(k: Int): Unit  
  {  
    s += 1  
  }  
}
```

```
1 to 2000
```



```
Predef.intWrapper(1).to(2000)
```

```
trait Function1[
  @specialized(scala.Int, scala.Long,
               scala.Float, scala.Double) -T1,
  @specialized(scala.Unit,
               scala.Boolean, scala.Int,
               scala.Float,
               scala.Long, scala.Double) +R]
  extends AnyRef
```

```
def foreach[@specialized(Unit) U] (f: Int => U)
```


Results in:

```
Predef.intWrapper(1).to(2000).
  foreach(new Function1[Int,Unit]
    { def apply(i: Int): Unit = {
      Predef.intWrapper(1).to(2000).
        foreach(new Function1[Int,Unit]
          { def apply(j: Int): Unit = {
            Predef.intWrapper(1).to(2000).
              foreach(new Function1[Int,Unit]
                { def apply(k: Int): Unit = {
                  s += 1
                })
              })
            })
          })
        })
      })
    })
```

Lots of overhead!

Ticket #1338: Optimize simple for loops

Tight loops - Scala

```
var s = 0
var i = 1; var j = 1; var k = 1

while (i <= 2000) { j = 1
  while (j <= 2000) { k = 1
    while (k <= 2000) { k += 1 }
    j += 1 }
  i += 1 }
```

60 μ s

For loop filtering

```
var a: Long = 0

for (i <- 1 to 2000;
     j <- 1 to 2000;
     k <- 1 to 2000
     if ((i+j+k) % 777 == 0)) a += 1
```

96 s

For loop filtering

```
var a: Long = 0

for (i <- 1 to 2000;
     j <- 1 to 2000;
     k <- 1 to 2000) {
  if ((i+j+k) % 777 == 0) a += 1
}
```

25 s

```
for (i <- 1 to 2000) if (i % 777 == 0) a += 1
```

Translates to:

```
Range(1, 2000).foreach(  
  i => if (i % 777) a += 1)
```

Further translates to:

```
Range(1, 2000).foreach(  
  new Function1[Int, Unit] {  
    def apply(i: Int): Unit = {  
      if (i % 777 == 0) a += 1  
    }  
  })
```

```
for (i <- 1 to 2000 if (i % 777 == 0)) a += 1
```

Translates to:

```
Range(1, 2000).withFilter(_ % 777 == 0)  
  .foreach(_ => a += 1)
```

```
def withFilter(p: A => Boolean):  
  FilterMonadic[A, Repr]
```

Further translates to:

```
Range(1,2000).withFilter(  
  new Function1[Object, Boolean] {  
    def apply(i: Object): Boolean = {  
      apply(Int.unbox(i)) }  
  
    def apply(i: Int): Boolean = {  
      i % 777 == 0 }  
  }) .foreach(  
    new Function1[Int, Unit] {  
      def apply(i: Int): Unit = { a += 1 }  
    })
```

Code spends 78% of time executing:

```
scala.runtime.BoxesRunTime.boxToInteger
```

Optimizing Higher-Order Functions in Scala

search for: **scala optimization**

by Iulian Dragos

Performance management

- Fine-grained optimization
- Efficiency of the overall design
- Algorithms choice
- Measure, use a profiler
- When in doubt, decompile and/or disassemble.

Performance management


Scala is a good fit for areas which have big performance needs.

Examples:

- server side components
- infrastructure
- concurrency

Performance-critical blocks of code

Using the right collection



mutable vs. immutable

performance

characteristics of different
collection types

	head	tail	apply	update	prepend	append	insert
immutable							
List	C	C	L	L	C	L	-
Stream	C	C	L	L	C	L	-
Vector	eC	eC	eC	eC	eC	eC	-
Stack	C	C	L	L	C	L	-
Queue	aC	aC	L	L	L	C	-
Range	C	C	C	-	-	-	-
String	C	L	C	L	L	L	-
mutable							
ArrayBuffer	C	L	C	C	L	aC	L
ListBuffer	C	L	L	L	C	C	L
StringBuilder	C	L	C	C	L	aC	L
MutableList	C	L	L	L	C	C	L
Queue	C	L	L	L	C	C	L
ArraySeq	C	L	C	C	-	-	-
Stack	C	L	L	L	C	L	L
ArrayStack	C	L	C	C	aC	L	L
Array	C	L	C	C	-	-	-

Table 24.10 · Performance characteristics of sequence types

	lookup	add	remove	min
immutable				
HashSet/HashMap	eC	eC	eC	L
TreeSet/TreeMap	Log	Log	Log	Log
BitSet	C	L	L	eC ^a
ListMap	L	L	L	L
mutable				
HashSet/HashMap	eC	eC	eC	L
WeakHashMap	eC	eC	eC	L
BitSet	C	aC	C	eC ^a

Table 24.11 · Performance characteristics of set and map types

Tail call optimization

```
def sum(x: Int): Int = {  
  if (x == 1) x  
  else x + sum(x - 1)  
}
```

770 μ s

```
sum(1000)
```

```
@tailrec
```

```
def sum(x: Int, total: Int = 0): Int = {  
  if (x == 0) total  
  else sum(x - 1, total + x)  
}
```

74 μ s

Type Specialization

```
class My [A] {  
  def iden(x: A) : A = x  
}
```

```
val a = new My[Int]  
for (i <- 1 to 1000000) a.iden(i)
```

40 ms

Type Specialization

```
class My [@specialized(Int) A] {  
  def iden(x: A): A = x  
}
```

```
val a = new My[Int]  
for (i <- 1 to 1000000) a.iden(i)
```

17 ms

Structural types

```
implicit def toInOperand[A] (a: A) = new {  
  def in(seq: Seq[A]) = seq.contains(a) }
```

```
def test {  
  var a: Long = 0  
  
  for (i <- 1 to 1000000) {  
    if (i in List(7,77,777)) a += 1  
    if (i.toString in  
        List("8","88","888")) a += 1  
    if ((i * 2) in  
        List (6, 66, 666)) a += 1  
  } }
```

455 ms

Structural types

```
class InOperand[A] (a: A) {  
  def in(seq: Seq[A]) = seq.contains(a)  
}
```

```
implicit def toInOperand[A] (a: A) = new  
InOperand(a)
```

```
def test {  
  var a: Long = 0  
  for (i <- 1 to 1000000) {  
    if (i in List(7,77,777)) a += 1  
    if (i.toString in  
        List("8", "88", "888")) a += 1  
    if ((i * 2) in List(6, 66, 666)) a += 1}}
```

403 ms

Structural types

```
implicit def toInOperand[A] (a: A) = new {  
    def in(seq: Seq[A]) = seq.contains(a) }
```

gets translated into:

```
def toInOperand[A] (a: A) : Object = {  
    new Object { def in(seq: Seq[A]) =  
        seq.contains(a) }  
}
```

Functional Programming

```
val r = (1 to 10000000).  
  map(_ * 2).  
  filter(_ > 1000000).  
  reduceLeft(_ + _)
```

96 ms

Functional Programming

```
val r = (1 to 10000000).view.  
  map(_ * 2).  
  filter(_ > 1000000).  
  reduceLeft(_ + _)
```

120 ms

for Range:

```
override def foreach[@specialized(Unit) U]  
(f: Int => U) {  
    if (length > 0) {  
        val last = this.last  
        var i = start  
        while (i != last) {  
            f(i)  
            i += step  
        }  
        f(i)  
    }  
}
```

```
def map [B] (f: (Int) => B) : Range[B]
```

implemented as:

```
for (x <- this) b += f(x)
```

which gets translated to:

```
this.foreach(x => b += f(x))
```

In case of `.view` code spends 25% of time in `scala.collection.immutable.Range.length`:

```
final def apply(idx: Int): Int = {  
    if (idx < 0 || idx >= length) throw new  
    IndexOutOfBoundsException(idx.toString)  
    start + idx * step}
```

```
def length: Int = fullLength.toInt
```

```
protected def fullLength: Long =  
    if (end > start == step > 0 && start != end)  
        ((last.toLong - start.toLong) /  
         step.toLong + 1) else 0
```

Functional Programming

```
var i = 1
var sum = 0

while (i <= 10000000) {
  if ((i * 2) > 1000000) sum += (i * 2)
  i += 1
}
```

48 ms

Traits vs. Abstract Classes

```
javap -c
```

```
...  
12: invokeinterface #26, 2; //InterfaceMethod  
Similarity.isNotSimilar:(Ljava/lang/Object;)Z  
...
```

```
...  
12: invokevirtual #24; //Method PointA.  
isNotSimilar:(Ljava/lang/Object;)Z  
...
```


Thank you!

References

- Odersky, Spoon, Venners: Programming in Scala, 2nd edition
- Wampler, Payne: Programming Scala
- <http://blog.danmachine.com/2011/01/moving-from-java-to-scala-one-year.html>
- <http://skillsmatter.com/podcast/scala/scala-performance-improvements-of-a-factor-of-4-to-30-boxing-and-specialization>
- <http://www.artima.com/underthehood/invocationP.html>
- <http://blog.richdougherty.com/2009/04/tail-calls-tailrec-and-trampolines.html>
- www.ibm.com/developerworks/java/library/j-jtp03253.html
- <http://scala-programming-language.1934581.n4.nabble.com/For-loops-dead-slow-in-Scala-td1942400.html>
- <http://stackoverflow.com/questions/2529736/groovy-scala-java-under-the-hood>
- <http://stackoverflow.com/questions/2794823/is-scala-functional-programming-slower-than-traditional-coding>
- <http://stackoverflow.com/questions/4001347/scala-compiler-says-my-method-is-recursive-in-case-when-implicits-and-anonymous-c/4002686#4002686>
- <http://stackoverflow.com/questions/4577144/which-scala-features-are-internally-implemented-using-reflection>
- <http://stackoverflow.com/questions/4962618/do-abstract-classes-in-scala-really-perform-better-than-traits>
- <http://stackoverflow.com/questions/4982529/how-to-recognize-boxing-unboxing-in-a-decompiled-scala-code>
- <http://stackoverflow.com/questions/504103/how-do-i-write-a-correct-micro-benchmark-in-java>